Explore: A Stock Discovery Experience for First-Time Investors

StockFox Senior Product Manager Assignment, Option 2. Ends at the Stock Health Scorecard.

Problem: New investors stall at three distinct points: no entry point into 5,000+ names, can't parse what surfaces, not enough trust to click.

Solution: Explore opens with brand recognition, not search. Four to five theme rails with disclosed sort criteria. Five-fact card that shows opinion-free facts and zeros scores/rankings.

Zero-candidate isn't one stall. It's three.

This submission designs a stock discovery experience for a first-time Indian retail investor with no first candidate in mind, where the terminal action opens the StockFox Scorecard. It is a conversation scaffold for a live two-way discussion, questioned line by line.

New investors stall at three points. Most products only fix one.

- **No entry point:** 5,000+ listed names, no obvious starting point. Each +10 options in a 401(k) plan drops enrollment ~1.5–2pp (Iyengar/Huberman/Jiang).
- **Can't parse:** ~4–7 item working-memory ceiling (Miller/Cowan). A user stalls on legible-but-overwhelming stock data.
- **Don't trust:** Overconfidence + disposition effect in Indian retail trading (~2.5M NSE investors, ~1.4B trades). 93% of F&O traders lost money FY22–24.
- **The gap:** Screener, Tickertape, Trendlyne, StockGro, TradingView all default new users to search, filters, or social feeds—none solve a zero-candidate start.
- **The signal:** StockGro reports ~90% of users exploring an idea on Google/social lose conviction before acting. Design targets that gap.

Explore opens with what a user already recognizes, not a search bar.

Zone A: Top bar

Slim search field, secondary position. Serves the tip-driven searcher who already has a name.

Zone B: Brand strip

Horizontal scroll of 8–10 recognizable consumer names (bank, airline, coffee chain). Entry point for zero-candidate users.

Zone C: Theme rails

4–5 horizontal rails, plain-language label + one disclosed objective sort criterion (e.g., "Profitable every year for 5 years"). 6–8 cards per rail.

Zone D: Compare tray

Persistent bottom tray, up to 4 stocks, added via +, entries stay in add-order, never re-sorted.

The card shows five facts and zero opinions.

Shown (5 facts): Name + description, sector tag, price + day change (small), rail's sort criterion, read-time/complexity tag.

Hidden (5 items): Composite Health Score, valuation calls, peer ranking, social signals, full 200+ metrics.

Why: A bare score on an unopened card reads as opinion, act now. The same score next to 200+ metrics reads as disclosed analysis. Not a formatting choice—it's regulatory: StockFox is not a SEBI Registered Analyst. SEBI's RA Regulations (amended Dec 2024) forbid price targets, buy/sell calls, and opinion on securities broadly.

The gap: Whether disclosed, objectively-sorted rails themselves count as a recommendation is unresolved. Flagged as an assumption, not a certainty.

Explore shows progress. It never rewards trading.

The mechanic: A single static readout: "You've explored 3 of 5 sectors this week." No streaks, points, badges, or push notifications.

- **FCA study:** Push notifications +11% trade count; points/prizes +12% trade count. Same mechanics: +8% risky-product trades (crypto/CFDs). Worst for low-literacy, female, younger users.
- **SEC:** Leaderboards, streaks, badges named a formal scrutiny category (RFI 2021-167).
- **Precedent:** $7.5M settlement with Robinhood; one customer made 12,700+ trades in 6 months under gamification.
- **Cost:** FCA's mechanics lift trade frequency ~11–12%. Explore will look less sticky week over week. That's the price of not reproducing the harm loop.

Browsing is free. Opening a Scorecard spends one of the three.

The rule: Explore (brand strip, rails, compare tray) is unmetered. The 3-free-analysis credit is spent only at the terminal action: opening a Scorecard.

- **Why here, not earlier:** Multi-session value threshold beats raw usage count as an activation predictor. Top freemium conversion flags: hitting a real limit, trying features beyond baseline—not browse volume.
- **Existing surfaces:** Confidence Journal (guided research), Learning layer (4-min lessons), Competitive Lens (peer comparison post-click) plug in here.
- **What doesn't plug in:** AI Research Copilot answers a query the user already knows how to ask. Explore serves the user with no query yet.
- **Cost:** Determined window-shoppers view dozens of cards free, zero funnel progress. Assumption: serving cost is small (cards are lightweight; Scorecard's AI analysis is expensive).

One metric: did a stranger open that stock's Scorecard.

Primary metric: Lift in Scorecard opens. Explore-sourced opens per weekly active Explorer vs. a search-only baseline cohort. Counts opens where entry point = Explore AND stock not already in watchlist AND view ≥10 seconds.

Guardrail 1 (harm): Confidence-score movement without Confidence Journal engagement. Flag any user whose score rises week over week while their journal steps = 0. Rising confidence with no research is overconfidence, not informed conviction.

Guardrail 2 (design integrity): Explore card bounce rate—Explore-sourced opens dismissed in <10 seconds. Catches clickbait optimization: if the primary metric rises because bounce rises too, the card got more sensational, not more useful.

Both guardrails assume Explore-specific instrumentation exists; verify with StockFox's data team before shipping.

Four decisions, four costs.

No composite score or ranking pre-click
Cost: No one-glance is-this-good shortcut; faster bounce for score-seeking users. Rejected: score with disclaimer. Reason: disclaimers don't survive contact with a number.

Meter only at Scorecard-open
Cost: Window-shoppers can browse indefinitely free. Nothing stops anxious users from postponing the click. Rejected: meter card views too. Reason: recreates the too-many-gates complaint.

No gamification tied to trading-adjacent actions
Cost: ~11–12% lower engagement lift than competitors. Rejected: scope streaks to Scorecard actions. Reason: FCA ties same mechanic to +6–8pp risky-product trade share.

Peer comparison hidden pre-click
Cost: Loses an early, useful comparative signal. Rejected: one-line peer delta on card. Reason: comparative claim reads closer to prefer-this-one.

What I don't know, and what I'd ask before this ships.

Five assumptions carry real risk if wrong.

A1:
No Explore-style discovery surface exists yet in the beta (confidence 0.4). Falsifier: Loom shows curated discovery already shipped.

A2:
The 3-free meter counts full Scorecard opens (confidence 0.5). Falsifier: Metered unit is something else (e.g., Copilot queries).

A3:
A pre-click card holds ~4–5 data points before tipping from inviting to intimidating (confidence 0.45). Falsifier: Usability testing shows novices want more.

A4:
Disclosed, objectively-sorted rails sit outside SEBI's opinion-on-securities definition (confidence 0.4). Falsifier: SEBI rules criteria-based grouping itself is a recommendation.

A5:
Confidence-score tracking is per-user and moves week to week (confidence 0.5). Falsifier: Score computed once; needs substitute signal.

Explore, zone by zone: interaction and empty-state detail.

- **First open, zero history:** Every user sees the same theme rails, non-personalized in v1. Cost: rails may not match individual interest. Traded for simplicity and not needing behavioral data before it exists.
- **Rails personalize only after signal:** At minimum one completed Scorecard open or an explicit tap of interest. Never before. Avoids cold-start recommender.
- **Loading state:** Tapping a card triggers a short loading state before the Scorecard renders, framed as pulling your analysis. Sets expectation that depth takes a moment.
- **Compare tray cap:** Tray caps at 4 stocks. A 5th add prompts a swap, not a scroll, keeping the shortlist within the same glance-level ceiling as the card itself.

Appendix A extends S2/S3 with interaction and empty-state detail for live defense.

The five named comparables, in full.

Product	Default Entry	Solves Zero-Candidate?
Screener.in	Search + 10,000+ community screens	No
Tickertape	Curated screens + sentiment gauge	No
Trendlyne	Screener-first + Red Flag tool	No
StockGro	Social feed + leaderboards + contests	No
TradingView	Screener + followable Ideas feed	No

All five default a new user to search, filter, or social—none lead with an opinionated starting point for a zero-candidate user. Confidence: 0.55 across the scan.

Three questions I'd bring to this conversation.

- **On the beta:** Does the current Loom beta already have a discovery/browse surface, and what did real users show about search-first vs. browse-first behavior? This is v1 against public surfaces only.
- **On regulation:** Is SEBI Research Analyst registration on the roadmap, or is staying unregistered deliberate long-term? That changes how conservative Explore needs to stay on opinions and comparisons.
- **On the funnel:** What does existing funnel data show for where users drop off before a first Scorecard open? Does it match no entry point as the biggest blocker, or is it something else?

Next step: Testing the ~4–5-item card and the rail sort criteria with 5–10 real novice users before any of it ships. This is v1, screen-level only.